

Akka vs. Google Gemini Enterprise Agent Platform

A comparison for teams building agentic AI

July 2026

The Gemini Enterprise Agent Platform is a bundle of pieces you assemble; Akka is one governed platform you run. At Cloud Next 2026, Google rebranded Vertex AI as the Gemini Enterprise Agent Platform — a code-first kit (ADK), a low-code builder (Agent Studio), and a managed runtime (Agent Engine) under a new name. The capability is real, but it is a collection of Google Cloud services the customer integrates, with the agent runtime excluded from the platform SLA, bound to Google Cloud for identity, models, and data. Akka delivers agents, memory, streaming, APIs, and governance as one platform — at six-nines availability, backed by indemnities — and Akka Specify can deliver and operate the entire governed system for you as a guaranteed outcome.

99.9999%

AKKA AVAILABILITY SLA

Weeks

AKKA SPECIFY DELIVERY

Fixed

ONE ALL-IN PRICE

90%

LESS INFRASTRUCTURE

DIMENSION	GEMINI ENTERPRISE AGENT PLATFORM	AKKA
What it is	A rebrand of Vertex AI (Cloud Next 2026): a set of Google Cloud services — ADK, Agent Studio, Agent Engine, Agent/Model Garden — for building agents	A full-stack agentic systems platform
Scope	Build tools plus a managed agent runtime; memory, governance enforcement, and the API/streaming tier are assembled from separate Google Cloud services	Orchestration, agents, memory, streaming, APIs, observability, and governance on one runtime
How you reach production	Self-assemble the platform's services, or contract a systems-integration engagement	Build with spec-driven development, or Akka Specify delivers and runs it
Commercial model	Pay-as-you-go meters plus per-seat subscriptions, plus integration labor	One fixed price — platform, infrastructure, tokens, training, delivery, and operations
Outcome guarantee	Effort only; the outcome is not guaranteed	The delivered outcome is guaranteed
Availability SLA	99.9% online inference (custom models on 2+ nodes); 99.5% Pipelines; the Agent Runtime is excluded from the SLA	99.9999% — entire platform, backed by indemnities
RTO / RPO	Per underlying Google Cloud service; customer-architected across regions	Sub-1-minute RTO; zero-byte RPO; active-active
Naming / churn	Renamed repeatedly: Enterprise Search → Generative AI App Builder → Vertex AI Search & Conversation → AI Applications → Agent Builder → Gemini Enterprise Agent Platform	Stable platform; 18 years; 100,000+ deployments
Governance / EU AI Act	Infrastructure governance: IAM, VPC-SC, CMEK, DLP, audit logs, data residency — no inline AI-policy enforcement, classification, or sealed audit artifact	Aspect-woven runtime enforcement + full pre-production governance
Lock-in	Google Cloud IAM identity, GCP-project data residency, models served through GCP	Akka cloud, any hyperscaler VPC, own Kubernetes, on-prem, or sovereign cloud; portable specs
Cost model	Pay-as-you-go: \$0.0864/vCPU-hour + \$0.0090/GB-hour runtime, \$0.25 per 1,000 stored memories, per-token model calls, per-query search	Shared compute; up to 90% lower infrastructure for the same workload; one fixed price

Improves after go-live	Not provided	Continuous evaluation, reinforcement learning, and distillation — up to 80% lower token cost over time
Certifications	Google Cloud compliance program	19 standards (SOC 2 II + public SOC 3, ISO 27001/42001, HIPAA, PCI DSS, GDPR, NIS2, DORA, EU AI Act, NIST AI RMF)

1. A Collection of Pieces, Not One Platform

The Gemini Enterprise Agent Platform is the new name for Vertex AI, announced at Google Cloud Next 2026 on April 22, 2026. It gives you three ways to build an agent — the code-first **Agent Development Kit (ADK)**, the low-code **Agent Studio** canvas, and prebuilt templates in **Agent Garden** — plus a managed **Agent Engine** runtime and a **Model Garden** of 200+ models (Gemini and Anthropic Claude among them). Each is a separate Google Cloud service the customer wires together and operates; the platform unifies the console and namespace, not the runtime.

Capability	Gemini Enterprise Agent Platform	Akka
Agent build tools	ADK + Agent Studio + Agent Garden	Built in
Managed agent runtime	Agent Engine — excluded from the SLA	Built in, covered by the platform SLA
Durable memory	Memory Bank, metered per stored memory	Built in, 4ms / sub-10ms
Real-time streaming	Provisioned from separate Google Cloud services	Built in, backpressured, petabyte-scale
Governance / policy enforcement	Infrastructure controls; no inline AI-policy enforcement	Inline, runtime-embedded
Pre-production governance	None	Classification, sign-offs, sealed posture

Akka delivers all of it as one runtime with one SLA. The customer does not integrate or operate the seams between an orchestration kit, a memory store, a streaming tier, and a governance stack.

2. Two Ways to Production: Assemble the Services, or Have Them Delivered

Reaching a production agentic system on the Gemini Enterprise Agent Platform means assembling ADK, Agent Studio, Agent Engine, Memory Bank, and separate streaming and governance services into one running system — yourself, or through a systems-integration engagement billed for effort. Akka offers a second path: **Akka Specify** takes your specifications and delivers and operates the entire governed system for you as a guaranteed outcome — in weeks, for one fixed price.

	With Gemini Enterprise Agent Platform	With Akka Specify
Model	Self-assemble the platform's services, or a systems-integration engagement	You provide the specifications
Who does the work	Your engineers, or a systems-integration engagement's staff	Akka generates, governs, delivers, and runs the system
Timeline	Months, by construction	Weeks, not quarters
Billing	Pay-as-you-go meters, per-seat subscriptions, plus integration labor	One fixed price
Afterward	You integrate, own, and operate every service and every seam	Kept up, safe, and improving — operated by Akka
Guarantee	Effort is guaranteed; the outcome is not	The delivered outcome is guaranteed

The Agent Runtime itself is explicitly excluded from the platform SLA — integrating these services into one running system, and operating them, is the customer's to assemble, whether in-house or through a systems-integration engagement. Akka Specify inverts that: you provide a handful of plain-language specifications, and Akka generates, tests, governs, deploys, and runs the system — then keeps it available, safe, and improving under one agreement, with one accountable owner.

3. Rebrand and Naming Churn

The product the customer would standardize on has been renamed repeatedly: Enterprise Search → Generative AI App Builder → Vertex AI Search & Conversation → AI Applications → Agent Builder → and now, at Cloud Next 2026, the Gemini Enterprise Agent Platform. The

rebrand is additive — existing Vertex AI workloads run unchanged under the new namespace, and SDKs, billing, and APIs were migrated with no breaking changes. But the churn is real and recent: the [vertexai.generative_models](#) Python SDK was deprecated June 24, 2025, with removal scheduled for June 24, 2026 — a live migration deadline. Akka has been one stable platform across 18 years and 100,000+ production deployments; the name on the contract has not changed under the customer.

4. Availability and Disaster Recovery

Google Cloud publishes a [99.9% monthly-uptime SLA](#) for online inference — and only for custom models deployed across 2 or more nodes, which the customer must architect and pay for. Vertex Pipelines carry 99.5% and the training-cluster control plane 99%. Most consequentially for agents: **the Agent Runtime (Agent Engine) is explicitly excluded from the SLA**, as are user-defined agents created in the Gemini Enterprise environment. Akka publishes a 99.9999% availability SLA across the entire platform — agents, memory, streaming, and governance — with sub-1-minute RTO, zero-byte RPO, active-active across regions, and contractual indemnities, operated by Akka's own SREs.

Metric	Gemini Enterprise Agent Platform	Akka
Availability SLA	99.9% online inference (2+ nodes)	99.9999%
Allowed downtime / year	~8.8 hours	~31 seconds
Agent runtime SLA	Excluded from the SLA	Covered
RTO / RPO	Per underlying service; customer-architected	Sub-1-minute / zero-byte
SLA scope	Individual Google Cloud services	The entire platform
Who operates it	The customer (or a systems-integration engagement)	Akka SREs, 24/7

5. Up to 90% Cheaper to Operate — and It Keeps Getting Cheaper

AI systems built with Akka are up to **90% cheaper to operate** than Python-based systems — a function of the infrastructure required for the same agentic transaction volume, not list price. The Gemini Enterprise Agent Platform bills **pay-as-you-go** on meters that scale with load: the Agent Engine runtime at \$0.0864 per vCPU-hour and \$0.0090 per GB-hour, stored sessions and memories at \$0.25 per 1,000, model calls per token, and search at \$1.50 per 1,000 queries — with per-seat subscriptions on the agent app layer on top.

Akka runs orchestration, agents, memory, streaming, APIs, observability, and governance on one shared-compute runtime. The efficiency comes from actor concurrency (~10T tokens/core/year vs ~2T; ~80% less compute than Python frameworks; Manulife: up to 300% more concurrency, 30–50% faster processing after porting from Python), shared compute, and micro-checkpointing. The spend is a predictable one fixed price, not meters that move with usage.

And the cost falls after go-live

The delivered outcome compounds. Akka Verify runs continuous evaluation, reinforcement learning on production and synthetic data, and distillation to smaller specialized models — cutting token cost up to 80% while raising accuracy over time. The Gemini Enterprise Agent Platform's pay-as-you-go meters do not fall on their own as usage matures; with Akka Specify, that tuning is part of the operated outcome, not a project the customer runs later.

7. Governance and the EU AI Act

The Gemini Enterprise Agent Platform provides strong infrastructure governance: Google Cloud IAM, VPC Service Controls, Customer-Managed Encryption Keys, Data Loss Prevention, Access Transparency, audit logging, and data residency. These secure the perimeter and the data. They do not enforce AI policy inline: no real-time guardrail/policy/judge layer woven into the runtime, no decision explainability, no human pause/override of a running agent as a platform primitive, no immutable interaction ledger, no pre-deployment classification against a regulatory corpus, and no sealed audit artifact.

The penalties are enforceable now

Violation	Maximum Fine
Prohibited AI practices (Art. 5)	EUR 35M or 7% global turnover
High-risk obligations (Art. 9-15)	EUR 15M or 3% global turnover

High-risk AI carries a 10-year logging-retention obligation (Art. 72).

6. Google Cloud Lock-In

The platform is bound to Google Cloud. Identity runs through Google Cloud **IAM** (or a federated IdP into it); data residency is scoped to a Google Cloud project and region (DRZ); and the Model Garden's models are served through GCP. Portability across this stack is the customer's project.

Akka deploys on Akka cloud, any hyperscaler VPC (AWS / Azure / GCP), the customer's own Kubernetes, on-prem, or a sovereign cloud with country-isolated data, networking, and local SREs — from one portable spec. The platform you standardize on is not the platform you are confined to.

How Akka governs

At the runtime: inline guardrails, policies, LLMs-as-a-judge, and sanitizers; hash-chained immutable evidence; HITL/HOTL control; classification against 186 regulations and 877 controls (288 carrying a financial penalty) before a system ships; multi-persona sign-offs; a sealed Governance Posture Package; and Akka Verify proving conformance from the running system. Governance the customer would otherwise assemble around Google Cloud's infrastructure controls, Akka enforces inline.

8. One Certified System — Built, Governed, Delivered, and Run

Building on the Gemini Enterprise Agent Platform means engineers in ADK or builders in Agent Studio producing agent logic; governance is a set of infrastructure controls a separate team configures around them, after the fact. Akka runs two independent lifecycles on one platform, and **Akka Specify** lets each audience work in plain language at its own level of abstraction:

BUILD LIFECYCLE FUNCTIONAL CONTRACT

"Rank incoming ER patients by acuity and route the top three to a clinician."

Product · developers · ML engineers · domain experts


v1.4 · versioned · tested

Akka Specify

ONE CERTIFIED AI SERVICE

Built, governed, delivered & run

- Agents, tools, orchestration, memory, APIs, streaming, UI
- Guardrails, sanitizers, HITL/HOTL, evaluations, halts
- Delivered, deployed, and operated for you



The build lifecycle and the governance lifecycle are versioned and tested independently, by different audiences, and through Akka Specify the whole certified system is delivered and operated as a guaranteed outcome — a workflow and a delivery model the Gemini Enterprise Agent Platform has no equivalent for.

9. Real-Time Streaming at Petabyte Scale

The Gemini Enterprise Agent Platform has no streaming engine; real-time pipelines are provisioned from separate Google Cloud services and wired to the agents. Akka's streaming is built into the runtime — continuous, backpressured, **petabyte-scale, in-memory**, with no external broker — powering both agent feedback loops and high-throughput data processing (the engine behind Tubi's real-time hyper-personalization at 5 billion tokens per second).

10. For the Buyer: Maturity, Adoption, and Accountability

Google Cloud is durable; that is not the question. The question is the **product's** maturity and the integration burden it puts on the buyer.

Buyer concern	Gemini Enterprise Agent Platform	Akka
Product maturity	Rebranded at Cloud Next 2026 (April); GA core, with agent components actively evolving; the prior SDK is being removed June 2026	Stable platform; 18 years; 100,000+ deployments (52 banks)
Naming / churn	Renamed five times (Enterprise Search → Agent Builder → Gemini Enterprise Agent Platform)	One platform, one name under the customer
Scope of accountability	The buyer integrates and operates ADK, Agent Engine, memory, streaming, and governance as separate Google Cloud services	One platform, one SLA, 24/7 SRE — Akka owns the running system
Delivery & outcome	Self-assemble the services, or a systems-integration engagement; effort is guaranteed, the outcome is not	Akka Specify delivers and operates the system for one fixed price; the outcome is guaranteed
Availability commitment	99.9% inference; Agent Runtime excluded from the SLA	99.9999% across the platform, with indemnities
Portability / lock-in	Google Cloud IAM identity, GCP-project data, models served through GCP	Akka cloud, any hyperscaler VPC, own Kubernetes, on-prem, or sovereign cloud; portable specs
Certifications & audits	Google Cloud compliance program	19 standards — SOC 2 II + public SOC 3, ISO 27001/42001, HIPAA, PCI DSS, GDPR, NIS2, DORA, EU AI Act, NIST AI RMF — plus annual pen tests, SBOMs, 40+ policies (trust.akka.io)
Budget predictability	Pay-as-you-go meters that scale with load	One fixed price finance can forecast

The decision is scope, accountability, and portability: the Gemini Enterprise Agent Platform gives the buyer a strong set of Google Cloud building blocks to assemble inside Google Cloud; Akka gives the buyer one governed platform, with one SLA, that runs anywhere — and, through Akka Specify, delivers and operates the finished governed system as a guaranteed outcome.

Customers Running Agentic and Real-Time Systems on Akka

Manulife 2,000 developers across 100 projects on one governed platform	Tubi 5B tok/s real-time hyper-personalization engine	Swiggy 71ms order-assignment AI, ~50% latency reduction	John Deere 1,000+ tractor sensors turned into real-time insight	Verizon 750% order-processing capacity gain; 6s → 2.4s response
--	--	---	---	---

Common Questions

We don't have a team to integrate this — what are our options?

Two. Build it on the platform with spec-driven development, or have Akka Specify deliver and operate it for you. You provide plain-language specifications; Akka generates, governs, delivers, and runs the system — agents, memory, streaming, APIs, and governance — as a guaranteed outcome, for one fixed price. With the Gemini Enterprise Agent Platform, production is self-assembly of ADK, Agent Studio, Agent Engine, and Memory Bank, or a systems-integration engagement you then own and operate.

How is Akka Specify different from a systems-integration engagement to assemble the Gemini Enterprise Agent Platform into a system?

A systems-integration engagement sells effort — time-and-materials over months — and hands back a system you own and operate; the outcome is not guaranteed. Akka Specify sells the outcome: a governed system delivered in weeks for one fixed price, then kept up, safe, and improving by Akka. You own the specifications, not the operational burden.

We are already on Google Cloud and starting with the Gemini Enterprise Agent Platform. Why add Akka?

The platform gives you good building blocks — ADK, Agent Studio, Agent Engine, Model Garden. A production agentic system also needs durable memory, streaming, an API tier, and inline runtime governance unified under one SLA. On Google Cloud you integrate those from separate services and own the seams; the Agent Runtime itself is excluded from the SLA. Akka delivers them as one platform at 99.9999%, and runs in your Google Cloud VPC if you want to stay there.

Sources

Rebrand / Cloud Next 2026: cloud.google.com/blog/products/ai-machine-learning/introducing-gemini-enterprise-agent-platform — Vertex AI rebranded April 22, 2026; existing workloads run unchanged

Components: docs.cloud.google.com/gemini-enterprise-agent-platform/overview, [/agent-studio/design-agents](https://docs.cloud.google.com/gemini-enterprise-agent-platform/agent-studio/design-agents), [/build/runtime](https://docs.cloud.google.com/gemini-enterprise-agent-platform/build/runtime); cloud.google.com/model-garden — ADK, Agent Studio, Agent Engine, Agent/Model Garden; 200+ models incl. Gemini and Anthropic Claude

Availability SLA: cloud.google.com/vertex-ai/sla, [/vertex-ai/generative-ai/sla](https://cloud.google.com/vertex-ai/generative-ai/sla) — 99.9% online inference (custom models on 2+ nodes), 99.5% Pipelines, 99% training control plane

Agent runtime SLA exclusion: cloud.google.com/terms/gemini-enterprise/sla — SLA does not apply to user-defined agents or to agents via Agent Runtime on the Gemini Enterprise Agent Platform

Pricing (pay-as-you-go): cloud.google.com/vertex-ai/pricing — Agent Engine \$0.0864/vCPU-hour + \$0.0090/GB-hour; \$0.25 per 1,000 stored memories; Vertex AI Search \$1.50 per 1,000 queries

Naming churn / SDK deprecation: docs.cloud.google.com/vertex-ai/generative-ai/docs/release-notes — `vertexai.generative_models` deprecated June 24 2025, removal June 24 2026; prior names incl. Enterprise Search, Generative AI

Vertex AI was just renamed, not removed — doesn't that mean it is stable?

The rebrand is additive and existing workloads run unchanged. But the agent layer has been renamed five times in a few years, and the prior generative-models SDK is being removed in June 2026. That is the kind of churn enterprise buyers weigh when they standardize. Akka has run as one platform under the same name for 18 years and 100,000+ deployments.

Can we govern for the EU AI Act with the platform's security controls?

You get IAM, VPC-SC, CMEK, DLP, audit logs, and data residency — infrastructure governance. The EU AI Act also expects AI-policy enforcement inline to the runtime: immutable records witnessed as they happen, human override of running agents, pre-deployment classification, and a sealed audit artifact. Akka embeds all of this and classifies a system against 186 regulations and 877 controls before it ships.

Isn't pay-as-you-go cheaper than a platform fee?

Pay-as-you-go means several meters that scale with load — runtime vCPU/GB-hours, stored memories, per-token model calls, per-query search — plus per-seat subscriptions on the app layer. Akka's shared-compute model is up to 90% cheaper to operate for the same agentic transaction volume, on one fixed price finance can forecast.

App Builder, Agent Builder

Governance / lock-in: docs.cloud.google.com/gemini-enterprise-agent-platform/machine-learning/general/vpc-service-controls; [/build/runtime](https://docs.cloud.google.com/gemini-enterprise-agent-platform/build/runtime) — IAM, VPC-SC, CMEK, DLP, Access Transparency, audit logging, data residency (DRZ)

Akka trust center: trust.akka.io — 19 compliance standards; SOC 2 II + public SOC 3; annual pen tests, SBOMs, 40+ policies

Akka performance: akka.io/blog/go-slow-to-go-fast — Manulife up to 300% more concurrency, 30–50% faster; ~10T vs ~2T tokens/core; ~80% less compute than Python

Akka Specify (spec-driven delivery): akka.io / akka.io/llms.txt — you provide the specifications; Akka generates, tests, governs, delivers, and operates the system as a guaranteed outcome; one fixed price covering platform, infrastructure, tokens, training, delivery, and operations; delivered in weeks; continuous improvement via reinforcement learning and distillation (up to 80% lower token cost, higher accuracy)

Akka platform: 99.9999% availability, active-active HA/DR, sub-1 min RTO, zero byte RPO (contractual indemnities); 186 regulations / 877 controls / 288 with a financial penalty; 100,000+ deployments / 18 years; profitable; Dell Technologies Capital