Akka Service Tiers — Choosing Your On-Ramp to Production Agentic AI

A TCO-driven guide to Akka's five service tiers for enterprise buyers, architects, and channel partners

July 2026 | akka.io

99.9999%	<1 min	0 byte	Up to 90%
AVAILABILITY GUARANTEE	RTO (ACTIVE-ACTIVE HA/DR)	RPO (ZERO DATA LOSS)	LOWER INFRA COST

THE CORE IDEA

Akka offers five tiers on a deliberate on-ramp from first experiment to global production. Four are fully resilient, production-grade environments. One — **Sandbox** — is intentionally lower quality-of-service, designed as a low-friction way to build and evaluate Akka in your own cloud without committing to a production footprint. Every deal starts by answering two questions: *where do you want Akka to run, and what does the workload need?* Every tier can also be delivered and operated for you: with **Akka Specify**, you provide the specifications and Akka builds, deploys, and operates the finished system as a guaranteed outcome — in weeks, for one fixed price covering the platform, the infrastructure, every inference and training token, and 24/7 operations.

The Five Tiers at a Glance

TIER	WHERE IT RUNS	QUALITY OF SERVICE	BEST FOR	TCO HEADLINE
STARTER	Akka's public cloud with private data link back to your environment	**Fully resilient, 24/7 SRE, no-downtime rolling updates**	Running a real production workload without installing anything in your own cloud	**Monthly billing in arrears, no POs, no invoices — spend scales with actual usage**
SANDBOX	A single region in your own VPC	Non-resilient, 9x5 support, not SLA-backed	Building and evaluating Akka inside your own cloud before a production commitment	**Lowest possible cloud footprint — roughly one-quarter the platform-core overhead of production tiers**
DAY 2 OPS	Two regions in your VPC (1 development, 1 production)	**Fully resilient, 24/7 SRE, elastic scale-to-zero, live CVE patching, AI explainability**	Running real workloads that don't yet need cross-region failover	**Replaces 6–9 separately billed services (orchestration + memory + streaming + APIs + evals + guardrails + logging + governance + observability) on shared compute**
BUSINESS CONTINUITY	Three regions in your VPC (1 development + 2 in active-active HA)	**Fully resilient + active-active HA/DR across regions or clouds**	Workloads that cannot tolerate a regional outage; revenue-bearing systems; regulated data flows	**Managed HA/DR with sub-1 minute RTO and zero-byte RPO — no DIY failover infrastructure to build or operate**
SOVEREIGN CLOUD	Three regions in your VPC, country-isolated with local SREs	**Fully resilient + HA/DR + private federation plane + country-level isolation**	Regulated industries, data residency mandates, sovereign requirements (EU, UK, Singapore, Japan, Australia, China, and more)	**Inherit 19+ InfoSec certifications rather than building compliance from scratch — full feature parity with commercial tiers**

Two Ways to Reach Any Tier: Build It, or Have Akka Specify Deliver It

Two commercial models sit on top of the five tiers above. A customer can build and operate the system with **spec-driven development** — provisioning whichever tier fits the workload and running it with their own team — or have **Akka Specify** deliver and operate the entire system as a guaranteed outcome: the customer provides the specifications, and Akka generates, governs, deploys, and operates the system in weeks, for one fixed price.

DIMENSION	SELF-BUILD (SPEC-DRIVEN DEVELOPMENT)	AKKA SPECIFY (DELIVERED & OPERATED)
Model	The customer provides requirements and builds on the platform	The customer provides the specifications
Who does the work	The customer's engineering team	Akka generates, governs, delivers, and runs the system
Timeline	As fast as the customer's team can ship	Weeks, not quarters
Billing	Service-core license, plus the customer's own platform-core infrastructure, plus the customer's build-and-operate time	One fixed price — platform, infrastructure, every inference and training token, delivery, and operations
Improves after go-live	Static once deployed, unless the customer invests in retraining	Continuous evaluation, reinforcement learning, and distillation — up to 80% lower token cost with higher accuracy over time
Tier it runs on	Any of the five tiers, the customer's choice	Typically Day 2 Ops or above
Afterward	The customer owns ongoing operations	Kept up, safe, and improving — operated by Akka
Guarantee	The customer controls execution; the outcome depends on their team	The delivered outcome is guaranteed

Every tier at Day 2 Ops and above already includes the personnel an Akka Specify engagement draws on: TAM/FDE (Technical Account Manager / Field Delivery Engineer) time scales from quarter-time at Day 2 Ops, to half-time at Business Continuity, to two dedicated FDEs at Sovereign Cloud. Akka Specify does not add a separate team to the deal — it uses the one already built into the tier.

Weeks TIME TO A DELIVERED SYSTEM	**Fixed** ONE ALL-IN PRICE	**24/7** OPERATED BY AKKA	**Up to 80%** LOWER TOKEN COST OVER TIME

The Two Starting Paths

Every deal begins with the same branching question: **does the customer want Akka-hosted or their-cloud-hosted?** The answer routes to one of two on-ramps. Both progress into Day 2 Ops, Business Continuity, or Sovereign Cloud as needs grow.

PATH A — STARTER

AKKA-HOSTED IN AKKA'S PUBLIC CLOUD

A fully resilient Akka region running in Akka's cloud with a private data link back to your environment. Monthly billing in arrears. No POs, no invoices, no procurement cycle. You never install anything.

Best for: *Teams with a real workload and no appetite for a BYOC commitment yet. Lines of business that need to ship without waiting on central IT. Proofs-of-value where the success criterion is "did it run in production?"*

PATH B — SANDBOX

INSTALLED IN YOUR OWN VPC

A single Akka region installed in your own cloud, deliberately configured for low overhead. No HA/DR. 9x5 support. Not SLA-backed. Designed as a low-friction way to evaluate Akka in your own environment without committing to a production footprint.

Best for: *Engineering teams who want Akka in their cloud before committing to a production tier. Enterprises whose procurement process requires a successful in-house pilot before larger contracts. Architecture teams running fit-for-purpose evaluation.*

> *Sandbox is cheap because it is deliberately stripped of the resilience layer. It is not a production tier, and migration to a higher tier is a rebuild — not an in-place upgrade. The word "Sandbox" does this positioning work for you: it signals throwaway, experimental, not the destination.*
>
> — Positioning principle: set expectations at sign-up, not at upgrade time

Quality of Service — Sandbox vs. the Production Tiers

This is the most important conversation to have on every deal. Customers who misunderstand it buy Sandbox for production workloads and blame Akka when something fails. Sandbox and the production tiers are not versions of the same thing at different price points — they are different QoS tiers built on different infrastructure configurations.

DIMENSION	SANDBOX	DAY 2 OPS, BUSINESS CONTINUITY, SOVEREIGN
HA/DR	None	**Day 2 Ops: none. BC + Sovereign: active-active across regions**
Rolling Updates	None — platform downtime during updates	**No-downtime rolling updates**
SRE Coverage	9x5 monitoring	**24/7 SRE monitoring and management**
SLA	Not SLA-backed	**Contractual 99.9999% availability, backed by indemnities**
CVE Patching	Requires downtime	**Live CVE patching with zero downtime**
Platform Core Footprint	**~5 platform cores (lowest infra overhead)**	**~20–25 platform cores per region (varies by tier)**
Intended Workloads	Evaluation, learning, early development	**Production, customer-facing systems, anything that matters**
Upgrade Path	Rebuild on a higher tier	**In-place progression between Day 2 Ops, BC, and Sovereign**

Platform Cores vs. Service Cores — The Akka Transparency Model

Akka is the only agentic AI platform that separates and discloses infrastructure overhead up front. Every other platform either hides it inside an opaque price, or forces customers to stand up separate services (vector DB, memory, streaming, evaluation, guardrails, logging, governance) each with their own infrastructure bill. Akka breaks the two apart.

How Akka Prices Infrastructure

Every Akka region installed in your VPC runs on two kinds of cores. You see both before you commit — no sticker shock at contract time.

PLATFORM CORES

What they run: Akka's own runtime infrastructure — clustering, resilience, data sharding, zero-trust networking, traffic steering. The infrastructure that delivers the 99.9999% availability guarantee.

Who pays: You provision these in your own cloud as part of your standard cloud infrastructure bill. **Akka does not charge a license fee for platform cores.** They are disclosed up front so you see the full footprint before you sign.

SERVICE CORES

What they run: Your agentic workloads — the agents, workflows, endpoints, entities, and views you build on Akka. The capacity that runs your business logic.

Who pays: Akka license scales with service cores. You pay for what you run, nothing more. When you add service cores, your Akka bill and your cloud infrastructure bill both scale predictably.

> *Sandbox is engineered to minimize platform-core overhead — roughly one-quarter the footprint of a production tier. The cost reduction comes from the infrastructure configuration itself, not from a discount. You pay less because you run less. When you need production-grade resilience, you rebuild on a tier that includes the full platform-core footprint required to deliver it.*
>
> — The Akka engineering & pricing principle: transparent sizing, not hidden overhead

The Third Dimension: One Fixed Price

Platform cores and service cores describe what a self-built deployment costs to run. An Akka Specify engagement replaces both lines — plus every inference and training token, delivery of the system, and 24/7 operations — with one fixed price. Rather than reconciling a license bill, a cloud infrastructure bill, metered token consumption, and a separate build-and-operate labor cost, the customer sees a single number covering the platform, the infrastructure, the tokens, the training, and the operations.

Matching Your Workload to the Right Tier

A five-step decision tree. Answer each question in order; the tier emerges in under a minute.

1. Where do you want Akka to run? → Akka's public cloud: Starter | Your own VPC: *continue to step 2*

2. Are you running something real, or exploring? → Exploring / evaluating: Sandbox | Running something real: *continue to step 3*

3. Can your workload tolerate a regional outage? → Yes: Day 2 Ops | No: *continue to step 4*

4. Do you need country-level data isolation for regulatory or sovereign reasons? → No: Business Continuity | Yes: Sovereign Cloud

5. Confirm the upgrade path before closing. → Sandbox → higher tier = **rebuild**. Day 2 Ops → BC → Sovereign = **in-place progression**.

TCO Arguments, Tier by Tier

TIER	PRIMARY TCO ARGUMENT	WHAT YOU'RE AVOIDING
STARTER	**Time-to-value and procurement friction. You're buying a running workload, not infrastructure. No capex, no architecture project, no deployment timeline.**	Avoided: multi-month deployment cycles, central-IT gating, the cost of not starting
SANDBOX	**Lowest-friction evaluation in your own environment. Smallest possible cloud footprint for validating fit before a larger commitment.**	Avoided: failed production rollouts based on documentation-only decisions; opportunity cost of over-committing too early
DAY 2 OPS	**Shared compute model. One platform replaces 6–9 separately billed services (orchestration, memory, streaming, APIs, evaluation, guardrails, interaction logging, governance, observability). Up to 90% infrastructure cost reduction vs. stitched stacks.**	Avoided: service sprawl, maintenance windows, patching downtime, Arize-class eval-tool spend, integration cost across point solutions
BUSINESS CONTINUITY	**Managed HA/DR with contractual guarantees. Sub-1 minute RTO, zero-byte RPO, backed by indemnities. Eliminates the DIY failover infrastructure project entirely.**	Avoided: multi-year, multi-million-dollar DIY failover projects; regulatory exposure from total state loss on region failure
SOVEREIGN CLOUD	**Full feature parity in a country-isolated configuration. Inherit 19+ InfoSec certifications (EU AI Act, ISO 42001, SOC 2, Singapore Agent Framework) rather than building your own posture. Runtime-native interaction logs, eval checkpoints, guardrail verdicts, legal hold, retention, and evidence export make Akka the evidence of record.**	Avoided: compliance remediation cost; audit preparation from scratch; separate eval evidence plane; EU AI Act fines up to 7% of global annual turnover

Two Different Savings, Not One

The infrastructure reduction above is a one-time, static number. A second, compounding saving comes from the system improving after go-live — and the two should not be confused with each other.

STATIC — INFRASTRUCTURE

Up to 90% less infrastructure than a stitched stack, from shared compute and platform-core disclosure. What the system costs to run today.

COMPOUNDING — TOKEN COST (SMART)

Akka Verify runs continuous evaluation, reinforcement learning on production and synthetic data, and distillation to smaller specialized models — cutting token cost up to 80% while raising accuracy over time. What the system costs to run a year from now.

> *Where Akka Specify delivers and operates the system, this self-improvement is part of the operated outcome, not a project the customer runs later. And whichever tier the system runs on, an Akka Specify engagement folds the platform, the infrastructure, every inference and training token, delivery, and 24/7 operations into one fixed price — a single line item finance can forecast.*
>
> — Akka Specify: spec-driven delivery, one fixed price, guaranteed outcome

Seven Questions That Expose the Hidden TCO Gap

Whether you are evaluating Akka for the first time or comparing it against Azure AI Foundry, LangChain, Temporal, or a DIY stack, these seven questions surface the real cost picture most vendors hide.

1. **"How many separately billed services are you stitching together today, and what's the combined monthly spend?"**
   Exposes the 6–9 service sprawl pattern. A single opaque price from one vendor is usually 6+ bills underneath the brand.
2. **"What are you paying for agent memory today, and what's the read/write latency?"**
   Surfaces the ~200ms bolt-on memory penalty. Akka delivers <10ms read/write durable memory as part of the platform.
3. **"What's your cloud spend on observability, tracing, and logging for AI workloads?"**
   Verbose agent tracing causes observability bills to balloon. Akka includes OTEL export without per-GB ingest charges.
4. **"Are you budgeting for Arize, LangSmith, Phoenix, or another eval/observability product?"**
   Exposes duplicated evaluation, trace storage, and audit-reconciliation cost. Akka captures eval checkpoints, guardrail verdicts, causal traces, token cost, legal hold, retention, and evidence export in the runtime.

5. "When a regulator or customer asks why an agent made a decision, which system is the evidence of record?"

Post-hoc dashboards are not the runtime. Akka's non-sampled interaction log is the authoritative record of what happened.

6. "When you need to patch a CVE in your runtime, how much downtime does that cost you?"

Exposes the lack of live CVE patching in most platforms. Akka patches with zero downtime — no maintenance windows.

7. "Have you priced out what it would cost to build active-active HA/DR in-house — and to staff, deliver, and operate the whole system?"

Exposes the hidden cost of DIY failover, delivery, and operations. Most enterprises estimate multi-year, multi-million-dollar projects that never finish. Akka ships HA/DR in BC and Sovereign, and Akka Specify folds delivery and operations into one fixed price.

Common Questions — Straight Answers

Can I just run production workloads on Sandbox to save money?

Sandbox is stripped of the resilience layer. There is no HA/DR, no 24/7 SRE, no SLA, and platform updates require downtime. If your workload has any production characteristics — customers depending on it, data that matters, revenue tied to uptime — Day 2 Ops is the minimum. Sandbox exists to let you evaluate Akka at the lowest possible cost, not to run production on the cheap.

Why is there no in-place upgrade from Sandbox?

Sandbox runs on a fundamentally different infrastructure configuration — fewer platform cores, no resilience layer, no HA/DR primitives. Retrofitting all of that onto a running environment would either force the same footprint as a production tier (defeating the purpose) or create an unreliable upgrade path. The rebuild is a feature. Build in Sandbox, prove out your design, then stand up the production tier your workload actually needs.

Why am I paying for platform cores on top of the license?

You are not paying Akka for platform cores — they are part of your own cloud infrastructure bill, not a line item on your Akka license. Akka charges only for service cores, where your workloads actually run. Platform cores are disclosed up front so you see the full picture before you commit. Most platforms hide this inside one opaque price.

We don't have a team to build this — what are our options?

Two. Build it with spec-driven development on whichever tier fits your workload, or have Akka Specify deliver and operate it. You provide plain-language specifications; Akka generates, governs, delivers, and runs the system as a guaranteed outcome, for one fixed price.

Azure AI Foundry has one bill. Isn't that simpler than Akka?

Azure AI Foundry is 6–9 separately billed services underneath a single brand — Azure OpenAI, Cosmos DB, AI Search, Event Hubs, API Management, App Insights, Content Safety, Private Endpoints, and often a separate eval/observability product. Each has its own pricing model and its own cost traps. The full bill arrives at the end of the month, after consumption has happened. Akka runs orchestration, agents, memory, streaming, APIs, evaluation checkpoints, guardrails, interaction logging, and governance on one platform with shared compute — and discloses platform-core overhead before you commit. Simpler is not one bill; simpler is one platform.

We already use Arize for evals and AI observability. Do we still need it?

Keep Arize for legacy workloads if it is already embedded, but do not budget it as required for Akka. Akka captures non-sampled interaction logs, evaluation checkpoints, guardrail verdicts, LLM calls, tool calls, token cost, authority snapshots, causal lineage, legal holds, retention, and evidence exports as part of the runtime. The question is whether you want to pay for a second evidence plane and then reconcile it with the runtime of record.

Can I start with Starter and migrate to Day 2 Ops later?

Yes — Starter is designed as an on-ramp. The common pattern is to prove the business case on Starter for one or two quarters, then bring the workload in-house to Day 2 Ops once your architecture, security, and procurement teams have signed off. Akka's field team helps you migrate with continuity of your running workload.

Do I really need 24/7 SRE on day one?

Every tier above Sandbox includes 24/7 SRE because production-grade infrastructure without 24/7 coverage is a contradiction. If a customer-facing workload goes down at 2am, you need someone watching. If your workload genuinely doesn't need 24/7 coverage, it isn't a production workload yet — which makes Sandbox the right fit, and the low-overhead configuration reflects that.

How is Akka Specify different from hiring an integrator to build this?

A labor-led integration engagement sells effort — time-and-materials, typically months — and hands back a system you then own and operate; the outcome is not guaranteed. Akka Specify sells the outcome: a governed system delivered in weeks for one fixed price, then kept up, safe, and improving by Akka. You own the specifications, not the operational burden.

The One-Minute Summary

QUESTION	ANSWER

How many tiers?	Five. Four are production-grade (Starter, Day 2 Ops, Business Continuity, Sovereign). One is a low-overhead evaluation tier (Sandbox).
Where does Akka install?	Starter: Akka's cloud. Everything else: your own VPC, on any hyperscaler or on-prem Kubernetes.
What's the fastest on-ramp?	Starter (monthly billing, no POs) or Sandbox (lowest footprint in your own cloud)
What determines my cloud bill?	Platform cores (Akka infrastructure, customer-provisioned) + service cores (your workloads, Akka-licensed). Both disclosed up front.
What's the migration story?	Sandbox → higher tier is a rebuild. Day 2 Ops → BC → Sovereign is in-place progression.
What's included in every production tier?	24/7 SRE, no-downtime rolling updates, live CVE patching, runtime evals, guardrails, interaction logging, evidence export, contractual 99.9999% availability, indemnity-backed SLAs
Who's the workload for?	Sandbox: evaluation. Day 2 Ops: real workloads, no failover. BC: workloads that can't tolerate regional outage. Sovereign: regulated or data-residency requirements.
What does Akka charge for?	Service cores (workload capacity) and TAM/FDE time. Platform cores are customer-provisioned cloud infrastructure, not an Akka line item.
How do I get it delivered?	Build it yourself with spec-driven development on any of the five tiers, or have Akka Specify deliver and operate the whole system for one fixed price — platform, infrastructure, tokens, training, delivery, and operations.