

## Akka vs. Vercel AI SDK

A comparison for teams building agentic AI

July 2026

**The Vercel AI SDK gives you outstanding developer experience for building and shipping AI features — but you own production.** It is a TypeScript developer SDK with serverless hosting: no enterprise durable runtime, no active-active HA/DR with zero-byte RPO, no six-nines SLA on the running agent, and no embedded governance. Akka is the agentic systems platform that owns the running system end to end and guarantees it — and Akka Specify can deliver and run the entire governed system for you as a guaranteed outcome.

**99.9999%**

AKKA AVAILABILITY SLA

**Weeks**

AKKA SPECIFY DELIVERY

**Fixed**

ONE ALL-IN PRICE

**90%**

LESS INFRASTRUCTURE

| DIMENSION                | VERCEL AI SDK                                                                                                  | AKKA                                                                                                    |
|--------------------------|----------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| What it is               | A TypeScript developer SDK for building AI apps and agents, with serverless hosting (AI Cloud / Fluid Compute) | A full-stack agentic systems platform — build on it yourself, or have it delivered and operated for you |
| How you reach production | Self-build, or contract a labor-led integration engagement                                                     | Build with spec-driven development, or Akka Specify delivers and runs it                                |
| Commercial model         | Consumption-metered serverless compute (Active CPU) + provisioned memory, plus build-and-operate labor         | One fixed price — platform, infrastructure, tokens, training, delivery, and operations                  |
| Outcome guarantee        | Effort only; the outcome is not guaranteed                                                                     | The delivered outcome is guaranteed                                                                     |
| Scope                    | Model calls, tool loops, streaming, UI hooks, MCP — outstanding DX for building; you own production            | Orchestration, agents, memory, streaming, APIs, observability, and governance on one runtime            |
| Language                 | TypeScript / JavaScript only (React, Next.js, Vue, Svelte, Node.js)                                            | Java / Scala; deploy on Akka cloud, hyperscaler VPC, own Kubernetes, on-prem, sovereign cloud           |
| Availability SLA         | 99.99% on Enterprise, commercially reasonable efforts, excludes Excused Downtime                               | 99.9999% — entire platform, backed by indemnities                                                       |
| HA / DR                  | Serverless redundancy; no published active-active SLA, RPO, or RTO commitment                                  | Active-active HA/DR; sub-1-minute RTO; zero-byte RPO                                                    |
| Durable state            | Stateless serverless functions; durability via event-log replay into distributed storage                       | Durable in-memory, 4ms reads / sub-10ms writes, replayable from its event journal                       |
| Governance / EU AI Act   | Security/privacy certifications; no AI-policy enforcement, classification, or sealed audit artifact in the SDK | Aspect-woven runtime enforcement + full pre-production governance                                       |
| Cost model               | Consumption-metered serverless compute (Active CPU) + provisioned memory, billed with load                     | Shared compute; up to 90% lower infrastructure for the same workload                                    |
| Improves after go-live   | Not provided                                                                                                   | Continuous evaluation, reinforcement learning, and distillation — up to 80% lower token cost over time  |
| Certifications           | SOC 2 Type II, ISO 27001:2022, PCI DSS, HIPAA, GDPR (access on request)                                        | 19 standards (SOC 2 II + public SOC 3, ISO 27001/42001, EU AI Act, NIST AI RMF)                         |

## 1. The Vercel AI SDK Is a Developer SDK with Hosting; Akka Is an Enterprise Platform

The Vercel AI SDK is the TypeScript toolkit for building AI-powered applications and agents — and it is excellent at that. It abstracts away the differences between model providers behind one clean API, and AI SDK 6 adds first-class agents, loop control, human-in-the-loop tool approval, full MCP support, and DevTools. Paired with Vercel's serverless hosting (AI Cloud and Fluid Compute), a team can build and ship an AI feature with remarkable speed.

That speed is real, and it is the right tool for building. The line is **production ownership**. The SDK is a library that runs inside functions you deploy; the durable runtime, the cross-region failover, the availability guarantee on the running agent, and the governance enforcement are not part of it. With the Vercel AI SDK you assemble and operate those concerns. With Akka they are the platform — and, through **Akka Specify**, Akka can build, deliver, and run the whole system for you.

| Capability                                    | Vercel AI SDK                                                                         | Akka                                                                     |
|-----------------------------------------------|---------------------------------------------------------------------------------------|--------------------------------------------------------------------------|
| Build AI features / agents in code            | Yes — best-in-class DX                                                                | Yes                                                                      |
| Model abstraction, tool loops, streaming, MCP | Yes                                                                                   | Yes                                                                      |
| Native durable runtime for the running agent  | Serverless functions + bolt-on durability                                             | Built in                                                                 |
| Active-active HA/DR (zero-byte RPO)           | Not published                                                                         | Built in                                                                 |
| Six-nines availability SLA on the agent       | 99.99% Enterprise                                                                     | 99.9999%                                                                 |
| Durable in-memory state                       | Distributed storage, replayed                                                         | Built in, 4ms / sub-10ms                                                 |
| Embedded runtime governance                   | Not in the SDK                                                                        | Inline, runtime-embedded                                                 |
| Pre-production governance                     | Not in the SDK                                                                        | Classification, sign-offs, sealed posture                                |
| Delivery & operation                          | The customer builds and operates the application, or contracts a labor-led integrator | Delivered and operated by Akka Specify as a guaranteed outcome, in weeks |

## 2. Two Ways to Production: Build It, or Have It Delivered

Akka offers two ways onto one platform: build the system yourself with spec-driven development, or have **Akka Specify** deliver and operate the entire governed system for you as a guaranteed outcome — in weeks, for one fixed price. The Vercel AI SDK offers neither a platform nor a delivery model: an application built with it is the customer's to build, deploy, and own.

|                   | With Vercel AI SDK                                                   | With Akka Specify                                      |
|-------------------|----------------------------------------------------------------------|--------------------------------------------------------|
| Model             | Self-build, or a labor-led integration engagement                    | You provide the specifications                         |
| Who does the work | Your engineers, or forward-deployed / staff-augmentation contractors | Akka generates, governs, delivers, and runs the system |
| Timeline          | 3–12 months, by construction                                         | Weeks, not quarters                                    |
| Billing           | Time-and-materials, plus consumption-metered serverless compute      | One fixed price                                        |
| Afterward         | You inherit the system and its operations                            | Kept up, safe, and improving — operated by Akka        |
| Guarantee         | Effort is guaranteed; the outcome is not                             | The delivered outcome is guaranteed                    |

Vercel's own Enterprise SLA covers the hosting platform, on a commercially reasonable efforts basis — the running agent, its durability, and its governance remain the customer's to build and operate. A team that cannot self-build contracts a labor-led integration engagement — billed for time and materials, guaranteeing effort rather than the outcome, and handing back a system the team then owns and operates. Akka Specify inverts that: you provide a handful of plain-language specifications, and Akka generates, tests, governs, deploys, and runs the system — then keeps it available, safe, and improving under one agreement.

## 3. Availability, Disaster Recovery, and Durable State

Vercel publishes a [99.99% availability SLA on its Enterprise plan](#), made on a commercially reasonable efforts basis and excluding Excused Downtime — outages caused by third-party vendors, the internet, your own software, or factors outside Vercel's reasonable control. Service credits are capped at 50% of the monthly fee, and the customer must file a claim with log files. There is no published active-active multi-region SLA, no zero-byte RPO, and no RTO commitment on the running agent.

Akka commits to **99.9999% availability** — the entire platform, contractually, backed by indemnities — with active-active HA/DR across regions, sub-1-minute RTO, and zero-byte RPO. The difference is roughly 52 minutes of allowed downtime per year versus about 31 seconds, and who owns — and operates — the recovery when a region fails.

| Metric                  | Vercel AI SDK                                        | Akka                        |
|-------------------------|------------------------------------------------------|-----------------------------|
| Availability SLA        | 99.99% (Enterprise, commercially reasonable efforts) | 99.9999%                    |
| Allowed downtime / year | ~52 minutes                                          | ~31 seconds                 |
| RTO                     | Not committed                                        | Sub-1 minute                |
| RPO                     | Not committed                                        | Zero byte                   |
| SLA scope               | Hosting platform; Excused Downtime excluded          | The entire running platform |
| Who operates it         | The customer                                         | Akka SREs, 24/7             |

State follows the same line. Vercel's compute is serverless — Fluid Compute functions are short-lived and reused, with a maximum duration of 800 seconds (1800s in beta); a function that exceeds it returns a 504 timeout. Durability for long-running agents comes from the Workflow DevKit, which records each step to an event log and replays it into isolated API routes, with production state in distributed storage. Akka holds state in **durable sharded in-memory** at 4ms reads and sub-10ms writes, replayable from its own event journal — no external store on the hot path.

## 4. Cheaper to Operate — and It Keeps Getting Cheaper

AI systems built with Akka are up to **90% cheaper to operate** than Python-based systems — a function of the infrastructure required to run the same agentic transaction volume, not list price. The drivers are actor concurrency (~10T tokens/core/year vs ~2T; ~80% less compute than Python frameworks), shared compute across orchestration, agents, memory, streaming, APIs, observability, and governance on one runtime, and micro-checkpointing that minimizes retries. Manulife reported up to 300% more concurrency and 30–50% faster processing after porting from Python.

Vercel's serverless model bills consumption: Active CPU while code executes, plus provisioned memory for running instances. Spend moves with load, and the meter covers the hosting compute — the agent runtime, memory, streaming, and governance you assemble around it carry their own cost. Akka runs all of it on one shared-compute runtime for one fixed price finance can forecast.

### And the cost falls after go-live

The delivered outcome compounds. Akka Verify runs continuous evaluation, reinforcement learning on production and synthetic data, and distillation to smaller specialized models — cutting token cost **up to 80%** while raising accuracy over time. The Vercel AI SDK provides no continuous evaluation, reinforcement learning, or distillation loop; improving a deployed agent is a project the customer builds and runs. With Akka Specify, that tuning is part of the operated outcome, not a project the customer runs later.

The spend is also predictable: **one fixed price finance can forecast** — covering platform, infrastructure, tokens, training, delivery, and operations — rather than consumption-metered serverless compute that moves with load, plus separate build-and-operate labor.

## 6. Governance and the EU AI Act

Vercel holds strong security and privacy certifications — SOC 2 Type II, ISO 27001:2022, PCI DSS, HIPAA, GDPR, and more — covering the hosting platform. The AI SDK does not provide AI-governance enforcement: no inline policy enforcement against regulations, no decision explainability, no human pause or override of a running process as a runtime guarantee, no immutable interaction ledger, no pre-deployment classification, and no sealed audit artifact. Those obligations are owned by the customer.

### The penalties are enforceable now

| Violation                         | Maximum Fine                         |
|-----------------------------------|--------------------------------------|
| Prohibited AI practices (Art. 5)  | <b>EUR 35M or 7% global turnover</b> |
| High-risk obligations (Art. 9-15) | <b>EUR 15M or 3% global turnover</b> |

High-risk AI carries a 10-year logging-retention obligation (Art. 72).

## 5. The Language Boundary

The Vercel AI SDK is **TypeScript / JavaScript only** — React, Next.js, Vue, Svelte, Node.js. That is the right surface for the web tier and for building AI features quickly. It also means an enterprise standardizing on it commits its agentic systems to the JS/TS runtime.

Akka runs on the JVM (Java / Scala) and deploys anywhere — Akka cloud, hyperscaler VPC (AWS/Azure/GCP), the customer's own Kubernetes, on-prem, or sovereign cloud with country-isolated data and local SREs. Specs are portable across those targets, so the deployment model is a procurement choice, not a constraint.

### How Akka governs

At the runtime: inline guardrails, policies, LLMs-as-a-judge, and sanitizers; hash-chained immutable evidence; HITL/HOTL control; classification against 186 regulations and 877 controls before a system ships; multi-persona sign-offs; a sealed Governance Posture Package; and Akka Verify proving conformance from the running system. Governance a Vercel AI SDK team would bolt on, Akka enforces inline.

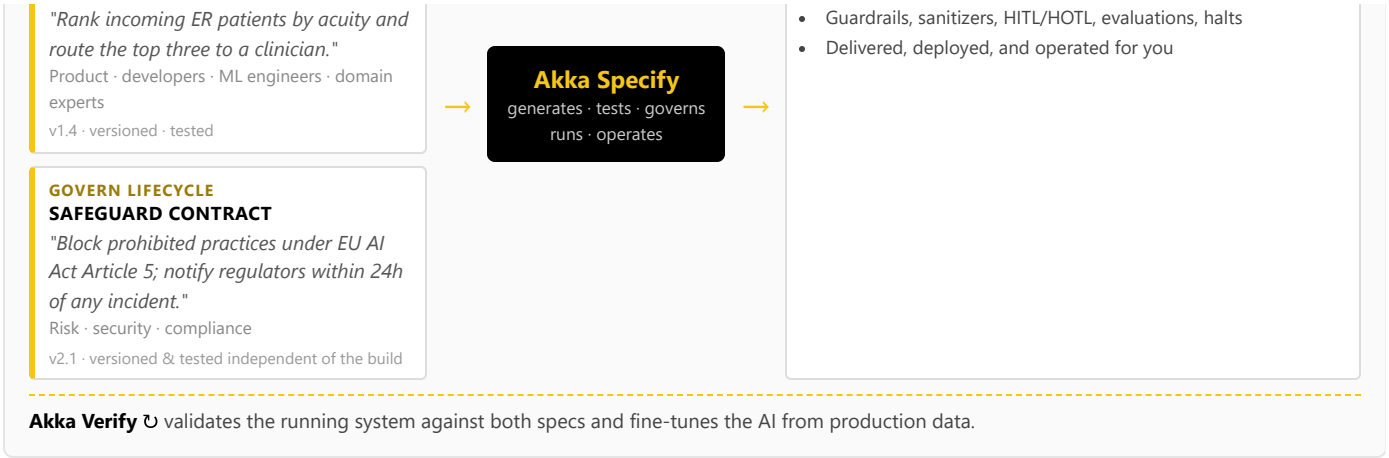
## 7. One Certified System — Built, Governed, Delivered, and Run

Building with the Vercel AI SDK means TypeScript engineers writing agent code; there is no built-in path for a product manager, domain expert, or risk officer to author and version a governance contract independently. Akka runs two independent lifecycles on one platform via **Akka Specify**:

**BUILD LIFECYCLE**  
**FUNCTIONAL CONTRACT**

**ONE CERTIFIED AI SERVICE**  
**Built, governed, delivered & run**

- Agents, tools, orchestration, memory, APIs, streaming, UI



Akka generates, tests, governs, **runs, and operates** one certified AI service that satisfies both specifications — agents, tools, orchestration, memory, APIs, streaming, and UI, together with guardrails, sanitizers, HITL/HOTL gates, evaluations, and interaction, evidence, and causal logging. The build lifecycle and the governance lifecycle are versioned and tested independently, by different audiences, and through **Akka Specify**, the whole certified system is delivered and operated as a guaranteed outcome — a workflow and a delivery model the Vercel AI SDK has no equivalent for.

## 8. Real-Time Streaming at Petabyte Scale

The Vercel AI SDK streams model tokens to the client, which is exactly what an interactive AI feature needs. It is not a streaming data engine. Akka's streaming is built into the runtime — continuous, backpressured, **petabyte-scale, in-memory**, with no external broker — powering both agent feedback loops and high-throughput data processing (the engine behind Tubi's real-time hyper-personalization at 5 billion tokens per second).

## 9. For the Buyer: Risk, Compliance, and Accountability

| Buyer concern                | Vercel AI SDK                                                                                                             | Akka                                                                                                                                                                           |
|------------------------------|---------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Certifications & audits      | SOC 2 Type II, ISO 27001:2022, PCI DSS, HIPAA, GDPR (access on request)                                                   | 19 standards — SOC 2 II + public SOC 3, ISO 27001/42001, HIPAA, PCI DSS, GDPR, NIS2, DORA, EU AI Act, NIST AI RMF — plus annual pen tests, SBOMs, 40+ policies (trust.akka.io) |
| Delivery & outcome           | Self-build or a time-and-materials integrator; effort is guaranteed, the outcome is not                                   | Akka Specify delivers and operates the system for one fixed price; the outcome is guaranteed                                                                                   |
| Scope of accountability      | The hosting platform; you build, integrate, and operate the agent runtime, memory, and governance                         | One platform, one SLA, 24/7 SRE — Akka owns the running system                                                                                                                 |
| Risk transfer                | Standard cloud terms; SLA credits capped at 50% of monthly fee                                                            | Availability and data-integrity guarantees backed by contractual indemnities                                                                                                   |
| Track record & funding model | Venture-funded: \$9.3B Series F (Sept 2025), ~\$863M raised; ~\$340M ARR run-rate (Mar 2026); profitability not disclosed | Profitable and self-funding; 18 years and 100,000+ deployments (52 banks); Dell Technologies Capital is largest shareholder, a customer, and an AI partner                     |
| Budget predictability        | Consumption-metered serverless compute that scales with load, plus build-and-operate labor                                | One fixed price finance can forecast                                                                                                                                           |

Vercel is well-capitalized and growing fast, and its developer experience is genuinely best-in-class. The decision is scope and accountability: the Vercel AI SDK gives developers an outstanding way to build and ship AI features; Akka gives you the enterprise platform that owns the running system — durable runtime, six-nines availability, active-active HA/DR, and embedded governance — and, through **Akka Specify**, will build, deliver, and run the whole system for you.

## Customers Running Agentic and Real-Time Systems on Akka

|                                                                                                  |                                                                                         |                                                                                            |                                                                                                    |                                                                                                    |
|--------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------|
| <p><b>Manulife</b></p> <p><b>2,000</b></p> <p>developers across 100 projects on one governed</p> | <p><b>Tubi</b></p> <p><b>5B tok/s</b></p> <p>real-time hyper-personalization engine</p> | <p><b>Swiggy</b></p> <p><b>71ms</b></p> <p>order-assignment AI, ~50% latency reduction</p> | <p><b>John Deere</b></p> <p><b>1,000+</b></p> <p>tractor sensors turned into real-time insight</p> | <p><b>Verizon</b></p> <p><b>750%</b></p> <p>order-processing capacity gain; 6s → 2.4s response</p> |
|--------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------|

## Common Questions

### We don't have a team to build this — what are our options?

Two. Build it on the platform with spec-driven development, or have Akka Specify deliver and operate it for you. You provide plain-language specifications; Akka generates, governs, delivers, and runs the system as a guaranteed outcome, for one fixed price. With the Vercel AI SDK, production is self-build or a labor-led integration engagement you then own.

### How is Akka Specify different from hiring an integrator to build our Vercel AI SDK app?

An integration engagement sells effort — time-and-materials over months — and hands back a system you own and operate; the outcome is not guaranteed. Akka Specify sells the outcome: a governed system delivered in weeks for one fixed price, then kept up, safe, and improving by Akka. You own the specifications, not the operational burden.

### We already use the Vercel AI SDK and love the developer experience. Why add Akka?

The Vercel AI SDK is an excellent way to build and ship AI features — keep using it for that. The question is what owns production: the durable runtime, cross-region failover, the availability guarantee on the running agent, and governance enforcement. Those are not in the SDK; you operate them yourself. Akka is the platform that owns the running system, with a 99.9999% SLA, active-active HA/DR, durable in-memory state, and embedded governance — or Akka Specify can deliver and operate the whole system for you.

## Sources

**Vercel AI SDK (TypeScript scope):** [ai-sdk.dev/docs/introduction](https://ai-sdk.dev/docs/introduction), [github.com/vercel/ai](https://github.com/vercel/ai) — “the TypeScript toolkit”; React, Next.js, Vue, Svelte, Node.js; open-source library, runs on any host

**AI SDK 6 (agents, loop control, MCP, DevTools, tool approval):**

[vercel.com/blog/ai-sdk-6](https://vercel.com/blog/ai-sdk-6), [ai-sdk.dev/docs/agents/overview](https://ai-sdk.dev/docs/agents/overview), [/agents/loop-control](https://ai-sdk.dev/docs/agents/loop-control)

**Vercel AI Cloud / Fluid Compute:** [vercel.com/blog/the-ai-cloud-a-unified-platform-for-ai-workloads](https://vercel.com/blog/the-ai-cloud-a-unified-platform-for-ai-workloads), [vercel.com/docs/fluid-compute](https://vercel.com/docs/fluid-compute), [/functions/limitations](https://vercel.com/docs/functions/limitations) — serverless; 800s max (1800s beta); 504 timeout

**Workflow DevKit / DurableAgent:** [vercel.com/blog/introducing-workflow](https://vercel.com/blog/introducing-workflow), [github.com/vercel/workflow](https://github.com/vercel/workflow) — event-log replay into isolated API routes; distributed storage in production

**Vercel Enterprise SLA:** [vercel.com/legal/sla](https://vercel.com/legal/sla), [vercel.com/enterprise](https://vercel.com/enterprise) — 99.99% commercially reasonable efforts; Excused Downtime excluded; credits capped at 50% of monthly fee

**Vercel certifications:** [vercel.com/docs/security/compliance](https://vercel.com/docs/security/compliance), [security.vercel.com](https://security.vercel.com) — SOC 2 Type II, ISO 27001:2022, PCI DSS, HIPAA, GDPR; access on request

**Vercel funding & revenue:** [sacra.com/c/vercel](https://sacra.com/c/vercel), [businesswire](https://businesswire.com) (Series F at \$9.3B,

### Doesn't Vercel's DurableAgent and Workflow DevKit give me a durable runtime?

The Workflow DevKit adds durability to TypeScript functions by recording each step to an event log and replaying it into isolated serverless routes, with production state in distributed storage. That makes individual workflows crash-safe. It is not an active-active HA/DR runtime with a zero-byte RPO and a six-nines SLA on the running system. Akka provides durable sharded in-memory state at 4ms/sub-10ms and a contractual 99.9999% availability guarantee.

### Can we add governance on top of the Vercel AI SDK?

You can add logging and evaluation libraries, but the EU AI Act expects enforcement inline to the runtime: immutable records witnessed as they happen, human override on running processes, and authorization captured at execution time. Bolt-on tools cannot gate a deployment or classify a system before it ships. Akka embeds inline guardrails, hash-chained evidence, and pre-deployment classification against 186 regulations and 877 controls, and proves conformance with Akka Verify.

### Isn't the Vercel AI SDK cheaper because it's open source?

The SDK is free and open source, and that lowers the cost of building. Production cost is the infrastructure to run the workload: Vercel bills consumption-metered serverless compute that scales with load, and the agent runtime, memory, streaming, and governance around it carry their own cost. Akka's shared-compute model is up to 90% cheaper to operate for the same agentic transaction volume, on one fixed price.

Sept 2025), [techcrunch.com](https://techcrunch.com) (Apr 2026) — ~\$863M raised; ~\$340M ARR run-rate; profitability not disclosed

**Akka Specify (spec-driven delivery):** [akka.io](https://akka.io) / [akka.io/llms.txt](https://akka.io/llms.txt) — you provide the specifications; Akka generates, tests, governs, delivers, and operates the system as a guaranteed outcome; one fixed price covering platform, infrastructure, tokens, training, delivery, and operations; delivered in weeks; continuous improvement via reinforcement learning and distillation (up to 80% lower token cost, higher accuracy)

**Akka trust center:** [trust.akka.io](https://trust.akka.io) — 19 compliance standards; SOC 2 II + public SOC 3; annual pen tests, SBOMs, 40+ policies

**Akka performance:** [akka.io/blog/go-slow-to-go-fast](https://akka.io/blog/go-slow-to-go-fast) — Manulife up to 300% more concurrency, 30–50% faster; ~10T vs ~2T tokens/core; ~80% less compute than Python

**Akka platform:** 99.9999% availability, active-active HA/DR, sub-1 min RTO, zero byte RPO (contractual indemnities); durable in-memory 4ms / sub-10ms; 186 regulations / 877 controls (288 with a financial penalty); 100,000+ deployments / 18 years; profitable; Dell Technologies Capital