

Akka vs. AWS Bedrock (AgentCore)

A comparison for teams building agentic AI

June 2026

AWS Bedrock AgentCore is a set of separate AWS services you wire and operate yourself — not an integrated, governed platform. AgentCore gives you seven modular building blocks (Runtime, Memory, Gateway, Identity, Code Interpreter, Browser, Observability) that you assemble on top of Bedrock model access, Guardrails, and Knowledge Bases — each billed and operated separately, all running inside AWS. The Akka Agentic AI Platform delivers orchestration, agents, memory, streaming, APIs, observability, and governance as one runtime that guarantees resilience and scalability, with a six-nines SLA Akka owns, deployable on any cloud.

99.9%

BEDROCK SLA

99.9999%

AKKA SLA

7

AGENTCORE SERVICES TO WIRE

Oct 2025

AGENTCORE GA

DIMENSION	AWS BEDROCK AGENTCORE	AKKA
What it is	A set of modular AWS services for building and hosting agents	The Akka Agentic AI Platform — a platform that guarantees resilience and scalability
Scope	Seven AgentCore modules assembled on Bedrock model access, Guardrails, and Knowledge Bases — the customer wires and operates the seams	Orchestration, agents, memory, streaming, APIs, observability, and governance on one runtime
Availability SLA	99.9% monthly for the Bedrock APIs, measured by request errors; service credits only	99.9999% — entire platform, backed by indemnities
RTO / RPO	Not published for the agent layer; inherits the AWS regional posture the customer architects	Sub-1-minute RTO; zero-byte RPO; active-active across regions
Deployment	AWS only — IAM, VPC/PrivateLink, CloudFormation, CloudWatch	Any cloud (AWS/Azure/GCP VPC), own Kubernetes, on-prem, sovereign cloud
Memory	AgentCore Memory — a separate, separately-billed service	Durable in-memory, 4ms reads / sub-10ms writes, built in
Governance / EU AI Act	Guardrails (model layer) + AgentCore Policy (tool layer), assembled; no embedded regulation/control enforcement	Aspect-woven runtime enforcement + pre-production classification against 190 regs / 1,040 controls
Cost model	Consumption-metered across the seven AgentCore services (per-vCPU-hour, per-request, per-token) plus separately billed model inference	Shared compute; up to 90% lower infrastructure for the same workload, fixed annual fee
Maturity	GA October 2025; Evaluations GA March 2026; Policy + Guardrails GA June 2026 — capabilities still landing	since 2007; 100,000+ production deployments; 52 banks
Certifications	Bedrock: ISO, SOC, CSA STAR L2, GDPR, FedRAMP High, HIPAA-eligible	19 standards (SOC 2 II + public SOC 3, ISO 27001/42001, EU AI Act, NIST AI RMF)

AgentCore Is a Set of Services to Assemble; Akka Is One Platform

AWS Bedrock AgentCore is a collection of separate AWS services the customer integrates and operates, not an integrated platform. At [general availability](#) it ships seven modular services — Runtime, Memory, Gateway, Identity, Code Interpreter, Browser, and Observability —

and an agent of substance also draws on Bedrock model access, Bedrock Guardrails, and Knowledge Bases. Each is provisioned, secured, billed, and operated on its own; the customer owns every seam between them.

Capability	AWS Bedrock AgentCore	Akka
Agent hosting / runtime	AgentCore Runtime (separate service)	Built in
Durable memory	AgentCore Memory (separate, separately billed)	Built in, 4ms / sub-10ms
Tool / API gateway	AgentCore Gateway (separate service)	Built in
Identity & authorization	AgentCore Identity (separate service)	Built in
Observability	AgentCore Observability on CloudWatch (separate)	Built in
Model access	Amazon Bedrock (separately billed per token)	Bring any model
Real-time streaming	Not a native AgentCore service	Built in, backpressured, petabyte-scale
Governance / policy enforcement	Policy + Guardrails, assembled	Inline, runtime-embedded

Akka delivers all of this on one runtime with one operational model. There are no seams to integrate, no per-service SLAs to reconcile, and no cross-service failure modes the customer inherits.

Availability and Disaster Recovery

The agent layer has no dedicated availability SLA. Amazon Bedrock publishes a [99.9% monthly uptime](#) commitment, measured as the percentage of Bedrock API requests that do not fail with errors, calculated per region, with service credits as the sole remedy. AgentCore inherits the regional, single-cloud posture the customer architects on top of AWS — the customer owns multi-region failover, RTO, and RPO for the assembled system.

Metric	AWS Bedrock AgentCore	Akka
Availability SLA	99.9% monthly (Bedrock API errors)	99.9999%
Allowed downtime / year	~8.8 hours	~31 seconds
RTO	Customer-architected	Sub-1 minute
RPO	Customer-architected	Zero byte
SLA scope	The Bedrock model APIs	The entire platform
Remedy	Service credits	Contractual indemnities

99.9% per region versus 99.9999% is the difference between roughly 8.8 hours and 31 seconds of allowed downtime a year. Akka's SLA covers the whole running system — orchestration, agents, memory, streaming, APIs, and governance — and Akka owns it with 24/7 SRE.

Up to 90% Cheaper to Operate

AI systems built with Akka are up to **90% cheaper to operate** than Python-based systems — a function of the infrastructure required for the same agentic transaction volume, not a list-price comparison. AgentCore meters its services independently — Runtime, Memory, Gateway, Identity, Code Interpreter, Browser, and Observability, plus Evaluations and Policy — on per-vCPU-hour, per-request, and per-token units, on top of standard Bedrock model token spend. Every layer the agent touches is a separate consumption meter that moves with load.

Akka runs all of it on one shared-compute runtime. The efficiency comes from actor concurrency (~10T tokens/core/year vs. ~2T; ~80% less compute than Python frameworks; Manulife: up to 300% more concurrency, 30–50% faster processing after porting from Python), shared compute, and micro-checkpointing that minimizes retries. The spend is a fixed annual fee — not a stack of usage meters that scale with traffic.

The EU AI Act Penalties Are Enforceable Now


What a violation costs

Violation	Maximum Fine
Prohibited AI practices (Art. 5)	EUR 35M or 7% global turnover
High-risk obligations (Art. 9-15)	EUR 15M or 3% global turnover
Incorrect information (Art. 99)	EUR 7.5M or 1.5% global turnover

High-risk AI carries a 10-year logging-retention obligation (Art. 72), enforceable since Feb 2025 (prohibited) / Aug 2025 (high-risk).

Two Lifecycles, One Certified System

Building on AgentCore means engineers wiring and operating a set of AWS services; there is no first-class path for a product manager, domain expert, or risk officer to contribute, and governance is assembled from separate services rather than running as its own lifecycle. Akka runs two independent lifecycles on one platform via **Akka Specify**:



Governance and the EU AI Act

AgentCore offers governance as assembled services, not embedded enforcement of regulation. [Bedrock Guardrails](#) operate at the model-inference layer (unsafe content, PII); AgentCore Policy operates at the tool-access layer. Both reached [GA in policy in June 2026](#). They enforce vendor-defined safety filters and tool rules — not the obligations of a named regulation. There is no built-in classification against the EU AI Act, no immutable hash-chained evidence ledger as a platform guarantee, no pre-deployment governance gate, and no sealed audit artifact.

How Akka governs

At the runtime: inline guardrails, policies, LLMs-as-a-judge, and sanitizers; hash-chained immutable evidence; HITL/HOTL control; atomic PII scrub-with-explain; pre-deployment classification against **190 regulations and 1,040 controls — 671 of which carry a financial penalty**; a multi-persona sign-off recipe engine; a sealed Governance Posture Package; and Akka Verify proving conformance from the running system. The governance lifecycle is versioned and tested independently of the build.

The build lifecycle and the governance lifecycle are versioned and tested independently, by different audiences — an audience and a workflow that a set of assembled cloud services has no equivalent for.

Real-Time Streaming at Petabyte Scale

Streaming is not a native AgentCore service; real-time pipelines are provisioned and operated separately on other AWS services. Akka's streaming is built into the runtime — continuous, backpressured, **petabyte-scale, in-memory**, with no external broker — powering both agent feedback loops and high-throughput data processing (the engine behind Tubi's real-time hyper-personalization at 5 billion tokens per second).

For the Buyer: Maturity, Lock-In, and Accountability

Buyer concern	AWS Bedrock AgentCore	Akka
Product maturity	GA October 13, 2025; core capabilities still landing (Evaluations GA March 2026; Policy + Guardrails GA June 2026)	since 2007, 100,000+ deployments, 52 banks
What you operate	Seven AgentCore services + model access + Guardrails + Knowledge Bases, integrated and run by the customer	One platform, one SLA, 24/7 SRE — Akka owns the running system
Portability / lock-in	AWS-only: IAM service-linked roles, VPC/PrivateLink, CloudFormation, CloudWatch, proprietary service APIs	Any cloud, own Kubernetes, on-prem, sovereign cloud; portable specs
Certifications & audits	Bedrock: ISO, SOC, CSA STAR L2, GDPR, FedRAMP High, HIPAA-eligible	19 standards — SOC 2 II + public SOC 3, ISO 27001/42001, EU AI Act, NIST AI RMF — plus annual pen tests, SBOMs, 40+ policies (trust.akka.io)
Risk transfer	Service credits on the model-API SLA	Availability and data-integrity guarantees backed by contractual indemnities
Budget predictability	Consumption meters across the AgentCore services plus model spend, scaling with load	Fixed annual fee finance can forecast

The structural point is not whether AWS will be around — it will. It is that AgentCore is a recent, consumption-metered collection of services the customer integrates and runs inside AWS, while Akka is an integrated, governed, portable platform with the SLA owned by the vendor.

Customers Running Agentic and Real-Time Systems on Akka

Manulife 2,000 developers across 100 projects on one governed platform	Tubi 5B tok/s real-time hyper-personalization engine	Swiggy 71ms order-assignment AI, ~50% latency reduction	John Deere 1,000+ tractor sensors turned into real-time insight	Verizon 750% order-processing capacity gain; 6s → 2.4s response
--	--	---	---	---

Common Questions

We're an AWS shop already. Why not just use AgentCore?

You can, but you are assembling and operating seven AgentCore services plus Bedrock model access, Guardrails, and Knowledge Bases — each billed and run separately, all locked to AWS. Akka delivers the same capabilities as one integrated runtime with a six-nines SLA Akka owns, and it runs in your AWS VPC, on another cloud, on your own Kubernetes, or on-prem. You get the platform instead of the integration project.

AgentCore is from AWS — isn't it the safe, mature choice?

AWS is durable; AgentCore is recent. It reached general availability in October 2025, and core governance and evaluation capabilities were still reaching GA into 2026 (Evaluations March 2026; Policy and Guardrails June 2026). Akka has operated since 2007 across 100,000+ production deployments behind a single, integrated platform.

Can't we add EU AI Act compliance with Guardrails and Policy?

Guardrails enforce content safety at the model layer and Policy enforces tool-access rules — both useful, neither is regulation enforcement. The EU AI Act expects classification before deployment, immutable records witnessed as decisions happen, human override of running processes, and a sealed audit artifact. Akka embeds classification against 190 regulations and 1,040 controls, hash-chained evidence, HITL/HOTL control, and pre-deployment governance inline.

Isn't AgentCore's consumption pricing cheaper than a fixed fee?

AgentCore meters its services on per-vCPU-hour, per-request, and per-token units, on top of model token spend, and the bill scales with load. Akka's shared-compute model is up to 90% cheaper to operate for the same agentic transaction volume, on a fixed annual fee finance can forecast.

Sources

AgentCore GA & components: aws.amazon.com/about-aws/whats-new/2025/10/amazon-bedrock-agentcore-available — GA October 13, 2025; seven modular services (Runtime, Memory, Gateway, Identity, Code Interpreter, Browser, Observability); VPC, PrivateLink, CloudFormation at GA

AgentCore preview: aws.amazon.com/about-aws/whats-new/2025/07/amazon-bedrock-agentcore-preview — announced in preview July 16, 2025

AgentCore maturity: aws.amazon.com/about-aws/whats-new/2026/03/agentcore-evaluations-generally-available (Evaluations GA March 2026); aws.amazon.com/about-aws/whats-new/2026/06/amazon-bedrock-agentcore-policy-guardrails-generally-available (Policy + Guardrails GA June 2026)

Bedrock SLA: aws.amazon.com/bedrock/sla — 99.9% monthly uptime, measured by Bedrock API request errors per region; service credits as remedy

AgentCore governance: aws.amazon.com/bedrock/guardrails (model-layer content/PII safety); Policy = tool-access-layer enforcement

AWS lock-in / VPC / IAM: docs.aws.amazon.com/bedrock-agentcore/latest/devguide/vpc-interface-endpoints.html; .../security-iam-

[awsmanpol.html](#) — IAM service-linked roles, VPC/PrivateLink, CloudWatch-backed observability

AgentCore pricing: aws.amazon.com/bedrock/agentcore — consumption-based across the AgentCore services (per-vCPU-hour + per-GB-hour runtime, plus memory, gateway, tools, observability); Bedrock model inference billed separately; exact rates per the AWS pricing page

Bedrock compliance: aws.amazon.com/bedrock/security-privacy-responsible-ai — ISO, SOC, CSA STAR L2, GDPR, FedRAMP High, HIPAA-eligible

Akka trust center: trust.akka.io — 19 compliance standards; SOC 2 II + public SOC 3; annual pen tests, SBOMs, 40+ policies

Akka performance: akka.io/blog/go-slow-to-go-fast — Manulife up to 300% more concurrency, 30–50% faster; ~10T vs ~2T tokens/core; ~80% less compute than Python

Akka platform: 99.9999% availability, active-active HA/DR, sub-1 min RTO, zero byte RPO (contractual indemnities); 190 regulations / 1,040 controls / 671 with a financial penalty; 100,000+ deployments / since 2007; profitable; Dell Technologies Capital